

```
In [398]: #import pandas library to be used  
import pandas as pd
```

Fall 2021 Data Science Intern Challenge

Assumptions:

Data set is as is and does not have any issues: No issues such as formatting, lack of headers and incorrect data types, headers were verified in the excel spreadsheet

All data in the imported set is valid for importation purposes.

All order amount values are in the same currency.

All shop_ids and user_ids have been entered correctly.

NOTE: Some tasks were performed outside of this python script, for information that was easily analysed from the google docs file, using excel formulas and sorting the data such as: Date range verification - all dates are in March of the same year.

Some resources used to solve this problem: <https://www.shopify.ca/blog/average-order-value>
(<https://www.shopify.ca/blog/average-order-value>) <https://www.optimizely.com/optimization-glossary/average-order-value/>
(<https://www.optimizely.com/optimization-glossary/average-order-value/>)
<https://www.bigcommerce.ca/ecommerce-answers/what-average-order-value/>
(<https://www.bigcommerce.ca/ecommerce-answers/what-average-order-value/>)

```
In [399]: #define the list of the dataset from the google docs file
data_source = "https://docs.google.com/spreadsheets/d/16i38oonuX1ylg7C_UAm
iK9GkY7cS-64DfiDMNiR41LM/export?format=csv&gid=0"

#import data as a csv
data_set = pd.read_csv(data_source)

print("This is the data set being used:\n", data_set, "\n")
```

This is the data set being used:

	order_id	shop_id	user_id	order_amount	total_items	payment_method
\						
0	1	53	746	224	2	cash
1	2	92	925	90	1	cash
2	3	44	861	144	1	cash
3	4	18	935	156	1	credit_card
4	5	18	883	156	1	credit_card
...
4995	4996	73	993	330	2	debit
4996	4997	48	789	234	2	cash
4997	4998	56	867	351	3	cash
4998	4999	60	825	354	2	credit_card
4999	5000	44	734	288	2	debit

	created_at
0	2017-03-13 12:36:56
1	2017-03-03 17:38:52
2	2017-03-14 4:23:56
3	2017-03-26 12:43:37
4	2017-03-01 4:35:11
...	...
4995	2017-03-30 13:47:17
4996	2017-03-16 20:36:16
4997	2017-03-19 5:42:42
4998	2017-03-16 14:51:18
4999	2017-03-18 15:48:18

[5000 rows x 7 columns]

The code above shows that there are 5000 rows of data and 7 columns, with the correct headers, verifying the data was imported correctly from the google docs file

```
In [400]: #given the somewhat large data set, convert from a dictionary to data frame
data_set = pd.DataFrame(data_set)
```

```
In [401]: #calculate an integer value for the number of rows (orders) in the data set to be used for future calculations
num_orders = len(data_set)

#calculate an integer value for the revenue
total_rev = sum(data_set.order_amount)
```

```
In [402]: AOV = total_rev / num_orders
          print("\nAOV =", AOV)
```

```
AOV = 3145.128
```

Analysing the Data

The average order value is verified to have been calculated correctly, with a rounded value of \$3145, as in the problem definition. With this, I can rule out the possibility of any arithmetic issues.

However, it is known that there is an issue with the data (from the problem statement), so first thing to do is ensure the data is valid by analysing the relevant columns. The shop_ids and the user_ids are correct, so is the created_at date

I will now check the following respectively:

1. order_id: To ensure that they are unique
2. price, using order_amount and total_items to calculate the unit price for analysis
3. quantity of items ordered

After this, I will recalculate the AOV and look at other values that can be used to measure the business growth, comparing them to the original data set and the corrected data set

1. Validating the order_ids

```
In [403]: #1. Validating that the order_ids are unique
          val_order_id_set = set(data_set.order_id)
          if len(val_order_id_set) != len(data_set.order_id):
              print("\nDuplicates exist\n")
          else: print("\nThere are not duplicates in the order_id\n")
```

```
There are not duplicates in the order_id
```

2. Analysing the price

```
In [404]: #2. To analyse the price, I will breakdown the order_amount to unit_price
          #This will tell me the price charged for a single shoe for each order
          #allowing a better understanding of the data set
```

```
unit_price = data_set.order_amount / data_set.total_items
```

```
In [405]: #update the data set to include the unit_cost
data_set['unit_price'] = unit_price
```

```
#Verify the data set
print("\nThis is the new data_set with the unit price of a shoe included:
\n",data_set)
```

This is the new data_set with the unit price of a shoe included:

	order_id	shop_id	user_id	order_amount	total_items	payment_method
\						
0	1	53	746	224	2	cash
1	2	92	925	90	1	cash
2	3	44	861	144	1	cash
3	4	18	935	156	1	credit_card
4	5	18	883	156	1	credit_card
...
4995	4996	73	993	330	2	debit
4996	4997	48	789	234	2	cash
4997	4998	56	867	351	3	cash
4998	4999	60	825	354	2	credit_card
4999	5000	44	734	288	2	debit

	created_at	unit_price
0	2017-03-13 12:36:56	112.0
1	2017-03-03 17:38:52	90.0
2	2017-03-14 4:23:56	144.0
3	2017-03-26 12:43:37	156.0
4	2017-03-01 4:35:11	156.0
...
4995	2017-03-30 13:47:17	165.0
4996	2017-03-16 20:36:16	117.0
4997	2017-03-19 5:42:42	117.0
4998	2017-03-16 14:51:18	177.0
4999	2017-03-18 15:48:18	144.0

[5000 rows x 8 columns]

The data set includes the unit price now.

To visualise the data better (instead of going through each row) I will plot the unit_price against shop_id


NOTE: In excel, I sorted by unit price and noticed there was an abnormally high unit price

This plot here is to visualise the data for better understanding

```
In [406]: #Each shop, since selling only one model of shoe charges the same price for that shoe
#NOTE: I chose to plot the unit_price against the shop_id because there are only 100 shops
#giving me an independent variable with fixed values.

data_set.plot(x='shop_id', y='unit_price', kind='scatter')
```

```
Out[406]: <AxesSubplot:xlabel='shop_id', ylabel='unit_price'>
```



From this plot, we can see that there is a single major outlier in the data set, where the price is about 25000 per unit (for a shoe).

I will now filter for the shop_ids that have the unit price >20000 (using the data from the plot) and group them by unit_price and shop_id

```
In [407]: unit_price_large = data_set.loc[data_set.unit_price > 20000]

unit_price_large_grouped = unit_price_large.groupby(["shop_id", "unit_price"])["shop_id"].count()
print(unit_price_large_grouped)


shop_id  unit_price
78        25725.0      46
Name: shop_id, dtype: int64
```

From the data set grouped by unit price and shop id, it shows that the unit_price of 25725 appears to be incorrect or a unique outlier. This also correlates with the shop ID 78 only, showing that only this shop has the issue in the data set, which I will exclude for purposes of accuracy, given how large it is compared to the other values.

NOTE: This step assumes that the unit_price of 25725 is either incorrect - from a bug, from incorrect data entry or a unique outlier. This high unit value would skew the results towards a higher AOV value, due to the relatively high revenue.

```
In [408]: #Correcting the data set to exclude the outlier price.
data_set_corrected = data_set.loc[data_set.shop_id != 78]
data_set_corrected.plot(x='shop_id', y='unit_price', kind='scatter')
```

```
Out[408]: <AxesSubplot:xlabel='shop_id', ylabel='unit_price'>
```



For further comparison, I will calculate the standard deviation for the original data set and compare it to the corrected data set.

```
In [409]: print('The original data unit_price has a standard deviation of:', data_set
            .std().unit_price, '\n'*2)
print('The corrected data unit_price has a standard deviation of:', data_set
      .std().unit_price, '\n'*2)
```

```
The original data unit_price has a standard deviation of: 2441.963725368454
6
```

```
The corrected data unit_price has a standard deviation of: 31.2602175328963
9
```

The corrected data set now appears more within reason

3. Analysing the quantity of shoes ordered

```
In [410]: total_items_grouped = data_set_corrected.groupby(["total_items"])["total_items"].count()
print("\nData set grouped by the total items, counting how many times the same quantity of shoes were purchased:\n\n",total_items_grouped)

total_items_grouped = pd.DataFrame(total_items_grouped)
#total_items_grouped = total_items_grouped.rename(columns = {total_items:IT}, inplace = False)
#print("\nData set grouped by the total items, counting how many times the same quantity of shoes were purchased:\n\n",total_items_grouped)

#total_items_grouped.plot(x='total_items', y='[total_items].count()', kind = 'scatter')
#print(len(total_items_grouped))
#total_items_groupedplt = data_set.groupby(["total_items"])["total_items"].plot(kind = 'scatter')
```

Data set grouped by the total items, counting how many times the same quantity of shoes were purchased:



```
total_items
1      1811
2      1816
3       932
4       292
5        77
6         8
8         1
2000     17
Name: total_items, dtype: int64
```

From this data, there appears to be a relatively large quantity of shoes (2000) ordered 17 times, when compared to the other number of orders ranging from 1 - 8, this data is clearly an outlier.

This may be a large retailer that buys a fixed amount of shoes with every purchase. Analysing the date data did not show a trend in the order dates, but showed that they ordered the same quantity each time. I will plot the total_items against the user id and shop id Note: I intuitively chose user_id and shop id here because the users are placing the orders and the shops are fulfilling them. I tested this plot with other values, such as unit_cost and decided the user_id and shop_id showed the data best

```
In [411]: data_set_corrected.plot(x='user_id', y='total_items', kind='scatter')
data_set_corrected.plot(x='shop_id', y='total_items', kind='scatter')
```

```
Out[411]: <AxesSubplot:xlabel='shop_id', ylabel='total_items'>
```



I will now create a dataset with the data where the total_items ordered is 2000, for further analysis, to see if there is a trend with the user_id or shop_id, or a relationship between both.


```
In [412]: #Dataset with high orders
total_items_large = data_set_corrected.loc[data_set_corrected.total_items
== 2000]

#grouped by user_id only
total_items_large_groupedu = total_items_large.groupby(["user_id"])[ "total_items"].count()

#grouped by shop_id only
total_items_large_groupeds = total_items_large.groupby(["shop_id"])[ "total_items"].count()

#grouped by both user_id and shop_id
total_items_large_groupedus = total_items_large.groupby(["user_id", "shop_id"])[ "total_items"].count()

print(total_items_large_groupedu, '\n')
print(total_items_large_groupeds, '\n')
print(total_items_large_groupedus, '\n')
```

```
user_id
607    17
Name: total_items, dtype: int64
```

```
shop_id
42     17
Name: total_items, dtype: int64
```

```
user_id  shop_id
607      42      17
Name: total_items, dtype: int64
```

Inspecting the data further in excel, looking at the total_items column, one thing to notice is the large total_items ordered come from a single shop ID 42, and user ID 607,

Given that all these orders were from a single store, it could be that these orders were made by a customer that regularly buys the same amount of shoes wholesale as previously mentioned

It is important to note that this shop sold at other quantities to other customers, and only user 607 ordered such a large amount each time.

Thus using this customer to calculate the AOV or other growth tracking values (such as the modal value) will skew the data higher. As such, I have decided to also exclude it from the data set.

I will now filter for where the total_items ordered is 2000 exclude this data, which is also where user_id is 607

```
In [413]: data_set_corrected2 = data_set_corrected.loc[data_set_corrected.user_id != 607]
print(data_set_corrected2)
```

	order_id	shop_id	user_id	order_amount	total_items	payment_method
\						
0	1	53	746	224	2	cash
1	2	92	925	90	1	cash
2	3	44	861	144	1	cash
3	4	18	935	156	1	credit_card
4	5	18	883	156	1	credit_card
...
4995	4996	73	993	330	2	debit
4996	4997	48	789	234	2	cash
4997	4998	56	867	351	3	cash
4998	4999	60	825	354	2	credit_card
4999	5000	44	734	288	2	debit

	created_at	unit_price
0	2017-03-13 12:36:56	112.0
1	2017-03-03 17:38:52	90.0
2	2017-03-14 4:23:56	144.0
3	2017-03-26 12:43:37	156.0
4	2017-03-01 4:35:11	156.0
...
4995	2017-03-30 13:47:17	165.0
4996	2017-03-16 20:36:16	117.0
4997	2017-03-19 5:42:42	117.0
4998	2017-03-16 14:51:18	177.0
4999	2017-03-18 15:48:18	144.0

[4937 rows x 8 columns]

I will now compare the standard deviation of the total_items from the previous data set to this data set

```
In [414]: print('In the previous data set (data_set_corrected) total_items has a standard deviation of:', data_set_corrected.std().total_items, '\n'*2)
print('In the corrected data set (data_set_corrected2) total_items has a standard deviation of:', data_set_corrected2.std().total_items, '\n'*2)
```

In the previous data set (data_set_corrected) total_items has a standard deviation of: 116.85728599309454

In the corrected data set (data_set_corrected2) total_items has a standard deviation of: 0.9828207382187315

The standard deviation is much less after the data is corrected, this shows the closeness in values of the total items after the large total items has been excluded

The AOV can now be calculated with this "incorrect" data excluded

```
In [415]: #recalculate the AOV
AOV_corrected2 = (sum(data_set_corrected2.iloc[:,3])) / len(data_set_corrected2)

print("""\nThe new AOV after the data has been cleaned up is:\n""", AOV_corrected2)
```

The new AOV after the data has been cleaned up is:
302.58051448247926

In summary, the issue with the data set appears to be abnormally high unit costs from a single shop, as well as abnormally high orders from a single customer.

Supplimental analysis

Below, is some analysis of the data sets at different levels show how the relationship between data changes as the data is cleaned up, using the mean, median and mode.

```

In [416]: mean_data_set = data_set.mean(axis = 0)
median_data_set = data_set.median()
mode_data_set = data_set.mode()

print("From the original data set:\n")
print(f""""The mean of the order amount and total item columns respective
ly:
Mean order amount:\n{mean_data_set.order_amount}
\nMean total items ordered:\n{mean_data_set.total_items}\n""")

print(f""""\nThe median of the order amount and total item columns respecti
vely:
\nMedian order amount:\n{median_data_set.order_amount}
\nMedian total items ordered:\n{median_data_set.total_items}\n""")

print(f""""\nThe mode of the order amount and total item columns respective
ly:
Mode order amount:\n{mode_data_set.order_amount[0]}
\nMode total items ordered:\n{mode_data_set.total_items[0]}\n""")

print('='*30, '\n'*2)

mean_data_set_corrected = data_set_corrected.mean(axis = 0)
median_data_set_corrected = data_set_corrected.median(axis = 0)
mode_data_set_corrected = data_set_corrected.mode(axis = 0)

print("From the initial data correction, excluding only the high priced sho
es of 25725\n")
print(f""""\nThe mean of the order amount and total item columns respective
ly:
Mean order amount:\n{mean_data_set_corrected.order_amount}
\nMean total items ordered:\n{mean_data_set_corrected.total_items}\n""")

print(f""""\nThe median of the order amount and total item columns respecti
vely:
\nMedian order amount:\n{median_data_set_corrected.order_amount}
\nMedian total items ordered:\n{median_data_set_corrected.total_items}\n
""")

print(f""""\nThe mode of the order amount and total item columns respective
ly:
Mode order amount:\n{mode_data_set_corrected.order_amount[0]}
\nMode total items ordered:\n{mode_data_set_corrected.total_items[0]}\n""
)

print('='*30, '\n'*2)
mean_data_set_corrected2 = data_set_corrected2.mean(axis = 0)
median_data_set_corrected2 = data_set_corrected2.median(axis = 0)
mode_data_set_corrected2 = data_set_corrected2.mode()

print("From the final data correction, excluding the high unit price of 25
725, and high number items orderd, 2000:\n")

print(f""""\nThe mean of the order amount and total item columns respective
ly:
Mean order amount:\n{mean_data_set_corrected2.order_amount}
\nMean total items ordered:\n{mean_data_set_corrected2.total_items}\n""")

print(f""""\nThe median of the order amount and total item columns respecti

```

```
vely:
\nMedian order amount:\n{median_data_set_corrected2.order_amount}
\nMedian total items ordered:\n{median_data_set_corrected2.total_items}\n
""")

print(f"""\nThe mode of the order amount and total item columns respective
ly:
Mode order amount:\n{mode_data_set_corrected2.order_amount[0]}
\nMode total items ordered:\n{mode_data_set_corrected2.total_items[0]}\n
""")
```

From the original data set:

The mean of the order amount and total item columns respectively:

Mean order amount:

3145.128

Mean total items ordered:

8.7872

The median of the order amount and total item columns respectively:

Median order amount:

284.0

Median total items ordered:

2.0

The mode of the order amount and total item columns respectively:

Mode order amount:

153.0

Mode total items ordered:

2.0

=====

From the initial data correction, excluding only the high priced shoes of 25
725

The mean of the order amount and total item columns respectively:

Mean order amount:

2717.3677836092047

Mean total items ordered:

8.851029471134437

The median of the order amount and total item columns respectively:

Median order amount:

284.0

Median total items ordered:

2.0

The mode of the order amount and total item columns respectively:

Mode order amount:

153.0

Mode total items ordered:

2.0

=====

From the final data correction, excluding the high unit price of 25725, and high number items ordered, 2000:

The mean of the order amount and total item columns respectively:

Mean order amount:

302.58051448247926

Mean total items ordered:

1.9947336439133077

The median of the order amount and total item columns respectively:

Median order amount:

284.0

Median total items ordered:

2.0

The mode of the order amount and total item columns respectively:

Mode order amount:

153.0

Mode total items ordered:

2.0

From the values above, the Mode and Median values remained the same regardless of the data set.

This is because the outliers were much less than the total data set, as such did not affect the mode and median

On the other hand, due to the way Mean is calculated, there is a vast difference in all 3 values

For this data set, using the median or mode, without correcting the data will give fairly accurate results, while the mean would not.